

Network security

Nicolò Fornari

February 27, 2016

Chapter 1

Network protocols

1.1 Introduction

Data link layer

It is the lowest logical level, the data link interconnects physical interfaces. Each interface is identified by a MAC address (Media Access Control).

The MAC address is 48 bit long, it is usually represented in Hex notation and it is used to route packets in local networks.

It uniquely identifies a network interface. It is assigned by the producer according to the standard IEEE 802.

Network Layer

IP operates at this level. IP addresses are dynamically assigned by an authority (eg. ISP's DHCP server).

Stateful: communication starts, develops,ends. eg. TCP

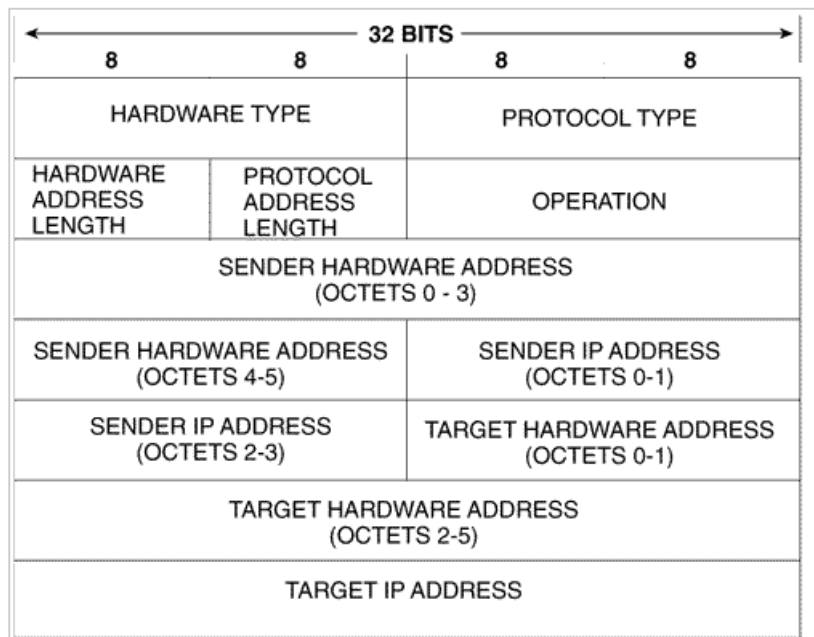
Stateless: IP

1.2 ARP

ARP (address resolution protocol) allows systems to associate an ip address to a MAC address. All addresses in the ARP table are added by one of these mechanisms:

- ARP request-reply:
who is 192.168.0.16 tell 192.168.0.1
192.168.0.16 is at 00-10-BC-2c-11-56
- Gratuitous ARP
192.168.0.16 is at 00-10-BC-2c-11-56

ARP frame header



ARP poisoning

The ARP protocol is declarative, it does not need an answer.

Nodes are not authenticated.

Limitations: it works only on LAN

Subnets and CIDR

Subnets are logical divisions of IP addresses. IP bits are partitioned as network,subnet,host.

A subnet mask indicates sections of IP addresses meant for network and subnet.

Eg. 255.255.255.0 means 24 bits for network and subnet and 8 bits for hosts.

CIDR

Classless Inter Domain Routing, it is a synthetic way to represent subnet masks.

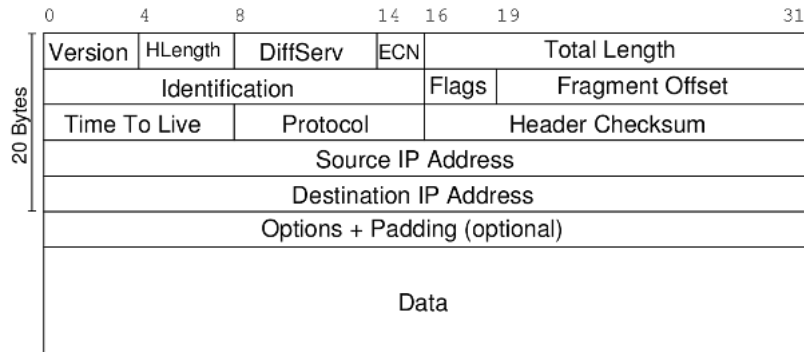
Example:

- Network mask: 255.255.0.0.
- CIDR representation: 132.132.1.10/16
- Hosts = 2^{16}

Formulas: (everything as binary)

- Network = Ip AND Subnet
- Host = Ip AND Not(Subnet)

1.3 IP



Some IPs are reserved for private networks:

- 10.0.0.0 → 10.255.255.255
- 192.168.1.1 → 192.168.255.255
- 172.16.0.0 → 172.16.255.255

Def A *datagram* is a basic transfer unit associated with a packet-switched network. The delivery, arrival time, and order of arrival need not be guaranteed by the network.


Def MTU maximum transmission unit

IP fragmentation *Identification*: 16 bit, is the unique identifier of the fragmented datagram. Note that all fragments have the same identification number.

Flags: 3 bits

- 0 Reserved, must be zero
- DF Don't fragment
 - If set to 0 → there may be fragments
 - If set to 1 → drop datagram if it has to be fragmented
- MF More fragments
 - 0 → last fragment
 - 1 → there are more fragments

Offset 13 bits, offset of this datagram wrt the first fragment with that ID

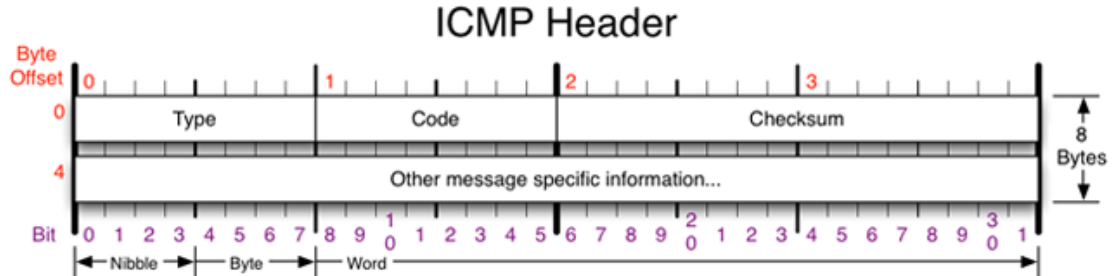
Fragmentation example

	A	B	C
Identification	4452	4452	4452
Flags	<ul style="list-style-type: none"> DF=0 MF=1 	<ul style="list-style-type: none"> DF=0 MF=1 	<ul style="list-style-type: none"> DF=0 MF=0
Offset	0	1480	2960

Remark 1. *DOS with IP fragments* You keep sending fragments without sending the first fragment, the router keeps waiting for it until it exhausts its memory.

1.4 ICMP

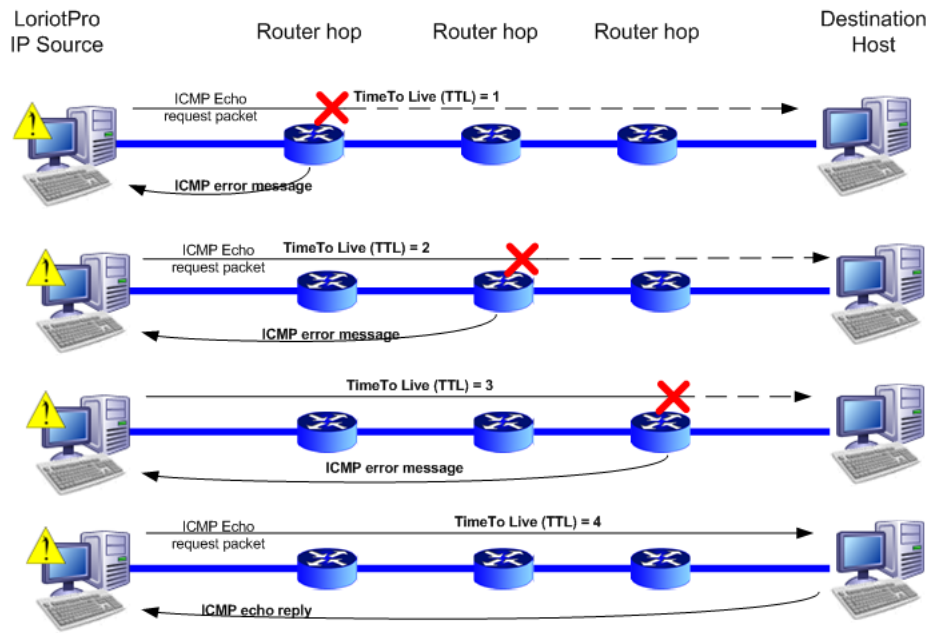
Internet control message protocol. It relies on IP and is an integral part of it.



Some message types

- 0 Echo Reply
- 3 Destination Unreachable
 - Code 0 → Net unreachable
 - Code 1 → Host unreachable
 - Code 2 → Protocol unreachable
 - Code 3 → Port unreachable
 - Code 4 → Fragmentation needed and DF set
 - Code 5 → Source route failed
- 4 Source Quench
- 5 Redirect
- 8 Echo
- 11 Time Exceeded
 - Code 0 → Net unreachable
 - Code 1 → Host unreachable
- 12 Parameter Problem
- 13 Timestamp
- 14 Timestamp Reply
- 15 Information Request
- 16 Information Reply

1.5 Traceroute



LUTEUS Copyrights 2008

1.6 Denial of service

Def a Denial of Service is a type of attack that aims at congesting or overpowering a system's capacity by generating requests the system will have to answer.

Examples

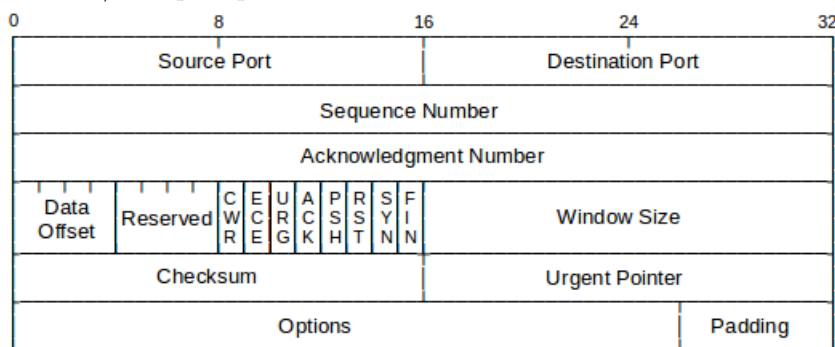
- Dos with IP fragmentation
- Ping Flooding (the attacker exploit his wider bandwidth)
- Ping of death

Chapter 2

Osi Transport Layer

2.1 TCP

The Transmission Control Protocol builds on top of IP the notion of state. Infact IP just delivers data while TCP manages the data segments by mean of checksums and re-delivery of unreceived/corrupted packets.



A server and a client that participate in a TCP connection open a *socket* which is a tuple (source-ip:source-port,destination-ip:destination-port).

Note that a client generates a source port randomly.



TCP details - some flags:

- **SYN:** initializes the TCP session, it should be set to 1 only for the first datagram
- **ACK:** acknowledge the reception of the segment
- **FIN:** it signals the intention of closing the connection
- **RST:** drop the connection (reset)

Sequence number: it is a 32 bit generated by each end (note that the sequence number of the client is independent from the server's one).

Acknowledgement number: 32 bits

2.2 Handshakes



2.3 Some TCP specifics

Both client and server set up a TCB (Transmission control block) to keep track of connection. The TCB structure is freed from memory when connection reaches status CLOSED.

A packet with RST flag up does not receive an answer.

If the state is CLOSED any packet with no RST receives a RST.

If the state is LISTEN

- SYN flag up + no ACK opens a TCP session. Answer is SYN+ACK
- Only ACK receives a RST
- Drop with no answer otherwise

2.4 SYN DoS

When the server receives SYN J it answers back with SYN K, ACK J+1.

The server opens a new session in a separate thread/ allocates resources. Then the server waits for the ACK K+1 from the client, it waits for the MSL (maximum segment lifetime set by default to 2 minutes). The same mechanism is on the sender side but of course the attacker controls it and can bypass it.

The attacker drops all SYN ACKs with a firewall to avoid exhausting its own resources. Note that a server typically has more bandwidth than a single client.

Remark 2. *The attack currently works in $\mathcal{O}(2n)$ because for each SYN a SYN ACK is received \rightarrow less bandwidth.*

However the attack can be improved by spoofing the source ip address, now the attacker is operating in $\mathcal{O}(n)$.

In theory the attack should not work because the server should receive a RST by each zombie, consequently free the TCB making the attack fail.

Still the attacker can choose a set of IPs which do not reply.

2.5 Dos mitigations

- Load balancing: distribute traffic loads evenly
- Rate limiter: deny traffic above a certain rate of SYN/sec
- Proof of work: require the source to solve a cryptopuzzle before allocating resources for the connection (note that this requires a protocol support).

2.6 TCP scans

It is possible to exploit specifications of a network protocol (TCP,UDP,...) to learn something about a system or a network.

SYN Scan: the attacker forges TCP packets with SYN=1. It is useful to see whether the remote system accepts incoming connections on a certain port.

Note that a polite way of scanning is performed in the following way: after the server's SYN ACK reply, the attacker sends a RST so that the 3-way handshake is never finished.

Host fingerprinting: different operating systems have their own independent implementation of the TCP stack.

Examples:

- FIN = 1
- all flags set to 0
- Xmas: FIN,URG,PSH = 1

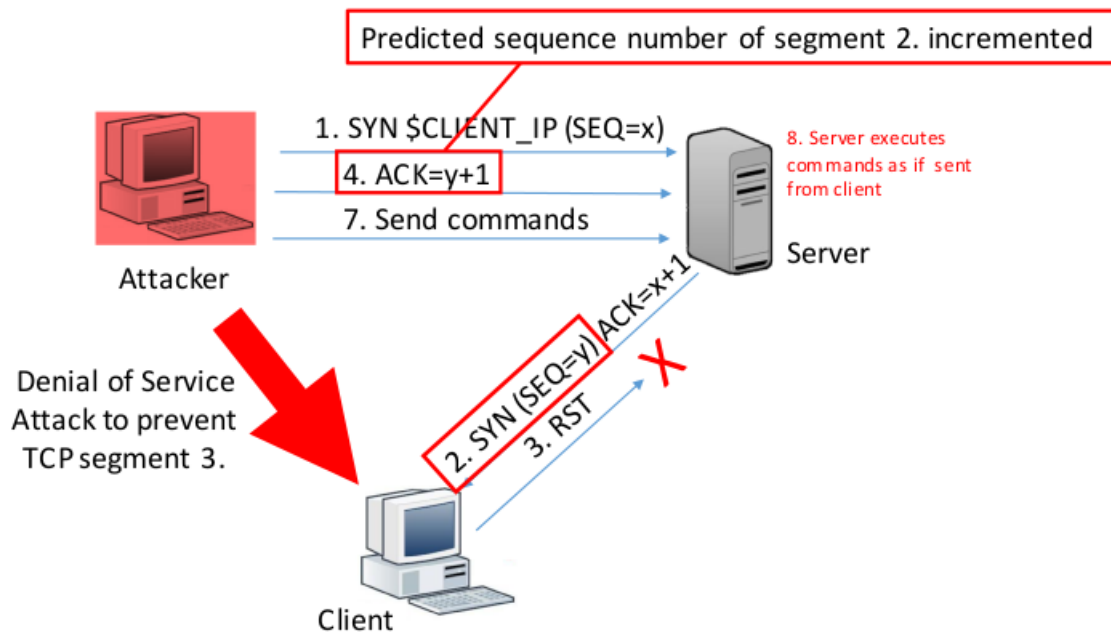
2.7 TCP Session hijacking - Mitnick attack

The attacker wants to send commands to a server they have no access to (eg. simple IP address authentication). The server has to think the attacker is the client, however the attacker does not sit in between client and server.

A TCP segment is identified and validated by:

- client ip \rightarrow known
- destination ip \rightarrow known
- port \rightarrow known (if not standard, scan)
- client SEQ number \rightarrow known (the attacker generates it)
- server SEQ number \rightarrow *unknown*

Back in the old days algorithms for generating the SEQ number were really simple.



2.8 UDP

The User Datagram protocol is a stateless. It is designed for fast delivery of data

- Data integrity can be controloed at application level
- It relies on the reliability of the underlying network link
- It does not guarantee delivery (there is no ACK mechanism)

Usage

- DNS servers
- NFS (network file systems)
- SNMP (simple network management protocol)
- DHCP (dinamic host configuration protocol)
- Most real time applications

2.9 UDP scans

It can be used to discover open ports on the network:

- CLOSE → ICMP port unreachable
- OPEN → no answer

However it is possible to configure a stealth system that does not reply to UDP requests to CLOSED ports by dropping ICMP packets with a firewall or router.

Chapter 3

Osi Session - Presentation - Application Layer

3.1 DNS

DNS (Domain Name Service) is a hierarchical system for domain name resolving. It translates human readable addresses to IP addresses the domain is reachable at. It uses UDP for fast answer (port 53).

Motivation: it is possible to assign multiple names to the same ip or viceversa. This flexibility is useful for:

- Server substitution
- Virtual hosting (a single server hosting multiple websites)
- A name is assigned to multiple IPs (so that the workload is distributed).

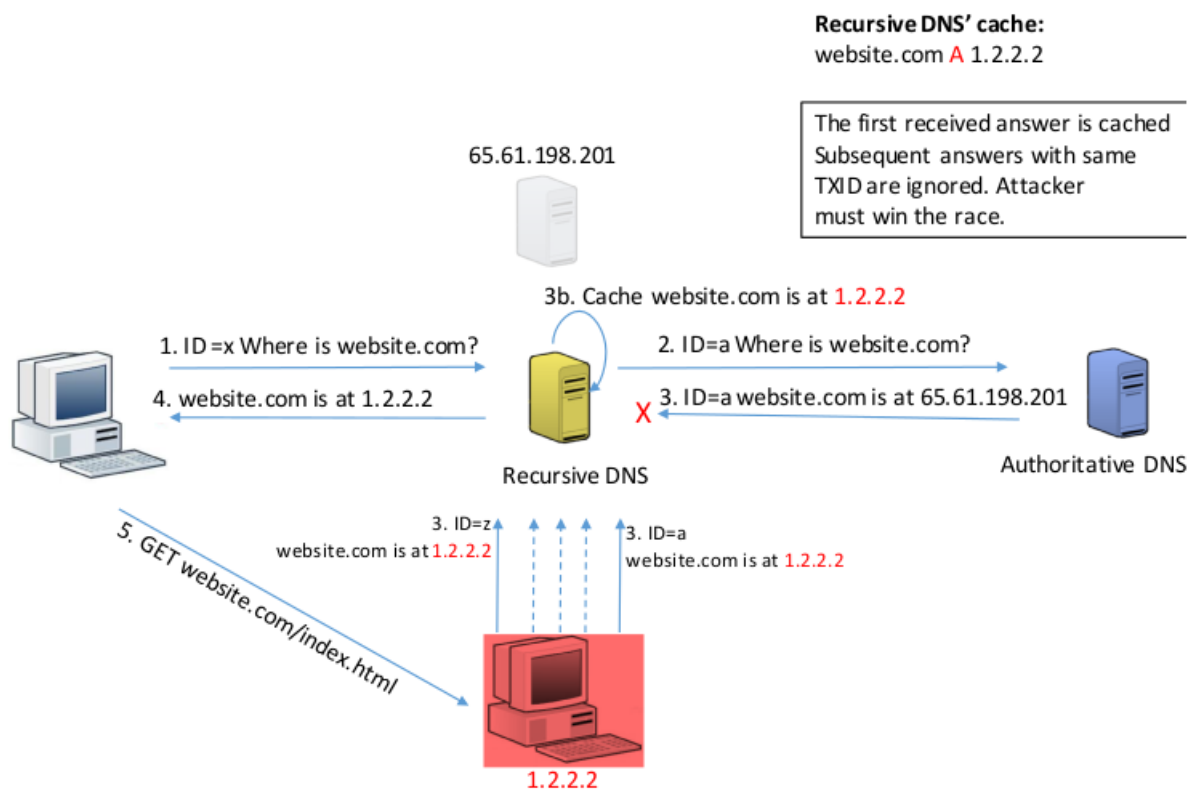
There exists different kind of DNS records, here we list a few:

- **A** correspondence name - IP(s)
- **AAAA** Same as A but works for ipv6
- **NS** ip of the DNS server to ask

DNS hierarchy

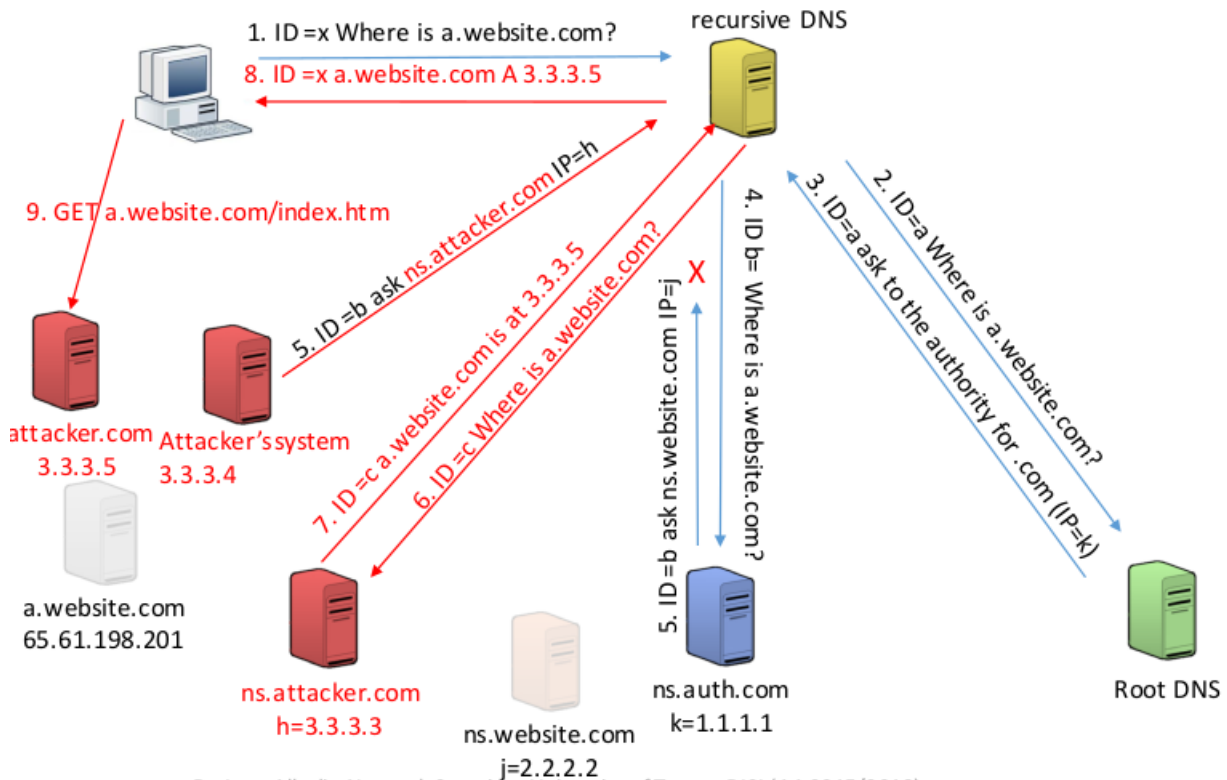
- Root DNSs → they are responsible for top level domain queries
- Authoritative DNS → A DNS that "knows the answer" (ie. it does not ask to other DNSs)
- Recursive DNS → It forwards queries to authoritative DNSs

3.1.1 DNS Cache poisoning



3.1.2 Kaminsky attack

The attacker rather than replacing an A record replaces an NS record, in this way he can get control over any (sub)domain.



Mitigation: the source of attack is low entropy with a 16 bit ID. Randomness is not enough and it is not feasible to change the protocol to 32 bits.

The solution is to randomize the source port to increase the entropy: any answer that does not match both source port and transaction ID will be dropped.

3.1.3 DNS amplification attack

It is DoS attack that exploits certain types of DNS answers that are much bigger in size than the requests.

Recall that the DNS works over UDP so the source IP is easy to spoof.

3.1.4 DNS zone transfer

A zone is a domain for which a server is authoritative. "Slave" servers can ask "authoritative" servers to copy their zone database (over TCP).

An attacker pretends to be a slave server and dump the zone DB. In this way he acquires knowledge of zone's infrastructure and it is eased in performing further attacks.

3.1.5 DNS Sec

It is a secure implementation of the DNS protocol: it implements DNS auth on top of normal DNS. Note that it protects just integrity and not confidentiality.

3.2 HTTP

HTTP is the main protocol on which the www works. It is based on the notion that a client can either request or submit data to a server. There are two methods: GET and POST.

HTTP is stateless, HTTP cookies enable statefulness.

3.2.1 Cookies

- Domain (who can read)
- Expires (if NULL valid only for a session)
- Secure (only over SSL)

HTTP session hijacking the attacker can read the session ID cookie and spoof the victim's identity (eg. facebook until 2011).

Secure cookies provide confidentiality but no integrity.

3.3 Telnet

It is a protocol used in remote control services. It operates over TCP port 23. Typically there is no authentication and no encryption.

3.4 Common issues

- Lack of authentication
- Communication channel is in the clear